

# Clustering

December 10, 2017

```
In [34]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

In [35]: teens = pd.read_csv('snsdata.csv')

In [36]: teens.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 40 columns):
gradyear      30000 non-null int64
gender        27276 non-null object
age           24914 non-null float64
friends       30000 non-null int64
basketball    30000 non-null int64
football      30000 non-null int64
soccer        30000 non-null int64
softball     30000 non-null int64
volleyball    30000 non-null int64
swimming      30000 non-null int64
cheerleading  30000 non-null int64
baseball      30000 non-null int64
tennis        30000 non-null int64
sports        30000 non-null int64
cute          30000 non-null int64
sex           30000 non-null int64
sexy          30000 non-null int64
hot           30000 non-null int64
kissed        30000 non-null int64
dance         30000 non-null int64
band          30000 non-null int64
marching      30000 non-null int64
music         30000 non-null int64
rock          30000 non-null int64
god           30000 non-null int64
```

```

church          30000 non-null int64
jesus           30000 non-null int64
bible           30000 non-null int64
hair            30000 non-null int64
dress           30000 non-null int64
blonde          30000 non-null int64
mall            30000 non-null int64
shopping        30000 non-null int64
clothes         30000 non-null int64
hollister       30000 non-null int64
abercrombie     30000 non-null int64
die             30000 non-null int64
death           30000 non-null int64
drunk           30000 non-null int64
drugs           30000 non-null int64
dtypes: float64(1), int64(38), object(1)
memory usage: 9.2+ MB

```

In [37]: teens.head()

```

Out[37]:   gradyear  gender    age  friends  basketball  football  soccer  softball  \
0      2006      M  18.982      7          0          0      0          0
1      2006      F  18.801      0          0          1      0          0
2      2006      M  18.335     69          0          1      0          0
3      2006      F  18.875      0          0          0      0          0
4      2006   NaN  18.995     10          0          0      0          0

   volleyball  swimming  ...  blonde  mall  shopping  clothes  hollister  \
0            0         0  ...      0    0          0         0          0
1            0         0  ...      0    1          0         0          0
2            0         0  ...      0    0          0         0          0
3            0         0  ...      0    0          0         0          0
4            0         0  ...      0    0          2         0          0

   abercrombie  die  death  drunk  drugs
0            0   0     0     0     0
1            0   0     0     0     0
2            0   0     1     0     0
3            0   0     0     0     0
4            0   0     0     1     1

[5 rows x 40 columns]

```

In [38]: teens.describe()

```

Out[38]:   gradyear    age  friends  basketball  football  \
count  30000.000000  24914.000000  30000.000000  30000.000000  30000.000000
mean    2007.500000    17.993950    30.179467     0.267333     0.252300

```

std	1.118053	7.858054	36.530877	0.804708	0.705357
min	2006.000000	3.086000	0.000000	0.000000	0.000000
25%	2006.750000	16.312000	3.000000	0.000000	0.000000
50%	2007.500000	17.287000	20.000000	0.000000	0.000000
75%	2008.250000	18.259000	44.000000	0.000000	0.000000
max	2009.000000	106.927000	830.000000	24.000000	15.000000

	soccer	softball	volleyball	swimming	cheerleading \
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	0.222767	0.161200	0.143133	0.13440	0.106633
std	0.917226	0.739707	0.639943	0.51699	0.514333
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	27.000000	17.000000	14.000000	31.000000	9.000000

	...	blonde	mall	shopping	clothes \
count	...	30000.000000	30000.000000	30000.000000	30000.000000
mean	...	0.098933	0.257367	0.353000	0.14850
std	...	1.942319	0.695758	0.724391	0.47264
min	...	0.000000	0.000000	0.000000	0.000000
25%	...	0.000000	0.000000	0.000000	0.000000
50%	...	0.000000	0.000000	0.000000	0.000000
75%	...	0.000000	0.000000	1.000000	0.000000
max	...	327.000000	12.000000	11.000000	8.000000

	hollister	abercrombie	die	death	drunk \
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	0.069867	0.051167	0.184100	0.114233	0.087967
std	0.346779	0.279555	0.624516	0.436796	0.399125
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	9.000000	8.000000	22.000000	14.000000	8.000000

	drugs
count	30000.000000
mean	0.060433
std	0.345522
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	16.000000

[8 rows x 39 columns]

```
In [39]: teens['age'].describe()
```

```
Out[39]: count      24914.000000
         mean        17.993950
         std         7.858054
         min         3.086000
         25%        16.312000
         50%        17.287000
         75%        18.259000
         max        106.927000
         Name: age, dtype: float64
```

```
In [40]: # Managing Outliers
```

```
def impute_age(cols):
    age = cols[0]
    if age >= 20:
        age = None
    else:
        if age < 13:
            age = None
        else:
            return age;
```

```
In [41]: teens['age'] = teens[['age']].apply(impute_age, axis = 1)
         teens['age'].describe()
```


```
Out[41]: count      24477.000000
         mean        17.252429
         std         1.157465
         min        13.027000
         25%        16.304000
         50%        17.265000
         75%        18.220000
         max        19.995000
         Name: age, dtype: float64
```

```
In [42]: teens['gender'].value_counts(dropna = False)
```

```
Out[42]: F      22054
         M       5222
         NaN     2724
         Name: gender, dtype: int64
```

```
In [43]: sns.countplot(x= 'gender', data=teens)
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f55842ec790>
```



```
In [44]: dummies= pd.get_dummies(data = teens['gender'])
         dummies.head()
```

```
Out[44]:
```

	F	M
0	0	1
1	1	0
2	0	1
3	1	0
4	0	0

```
In [45]: teens= pd.concat([teens, dummies], axis= 1)
         teens.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 42 columns):
gradyear      30000 non-null int64
gender        27276 non-null object
age           24477 non-null float64
friends       30000 non-null int64
basketball    30000 non-null int64
football      30000 non-null int64
soccer        30000 non-null int64
softball     30000 non-null int64
volleyball    30000 non-null int64
```

```

swimming      30000 non-null int64
cheerleading  30000 non-null int64
baseball      30000 non-null int64
tennis        30000 non-null int64
sports        30000 non-null int64
cute          30000 non-null int64
sex           30000 non-null int64
sexy          30000 non-null int64
hot           30000 non-null int64
kissed        30000 non-null int64
dance         30000 non-null int64
band          30000 non-null int64
marching      30000 non-null int64
music         30000 non-null int64
rock          30000 non-null int64
god           30000 non-null int64
church        30000 non-null int64
jesus         30000 non-null int64
bible         30000 non-null int64
hair          30000 non-null int64
dress         30000 non-null int64
blonde        30000 non-null int64
mall          30000 non-null int64
shopping      30000 non-null int64
clothes       30000 non-null int64
hollister     30000 non-null int64
abercrombie   30000 non-null int64
die           30000 non-null int64
death         30000 non-null int64
drunk         30000 non-null int64
drugs         30000 non-null int64
F             30000 non-null uint8
M             30000 non-null uint8
dtypes: float64(1), int64(38), object(1), uint8(2)
memory usage: 9.2+ MB

```

```
In [46]: teens[['gender', 'F', 'M']].head()
```

```

Out[46]:   gender  F  M
0         M  0  1
1         F  1  0
2         M  0  1
3         F  1  0
4        NaN  0  0


```

```

In [47]: plt.figure(figsize=(12, 7))
         sns.boxplot(x='gradyear', y='age', data=teens, palette='winter')

```

Out[47]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f5584227190>



```
In [48]: GradyearMeansByAge = teens['age'].groupby(teens['gradyear']).mean()
```

```
In [49]: print(GradyearMeansByAge)
```

```
gradyear
2006    18.655858
2007    17.706172
2008    16.767701
2009    15.819573
Name: age, dtype: float64
```

```
In [50]: def impute_age(cols):
          Age = cols[0]
          Gradyear = cols[1]

          if pd.isnull(Age):

              if Gradyear == 2006:
                  return 18.655858

              elif Gradyear == 2007:
                  return 17.706172
```

```

        elif Gradyear == 2008:
            return 16.767701

        else:
            return 15.819573

    else:
        return Age

```

```
In [51]: teens['age'] = teens[['age', 'gradyear']].apply(impute_age,axis=1)
```

```
In [52]: teens['age'].describe()
```

```

Out[52]: count      30000.000000
         mean         17.237326
         std          1.141821
         min          13.027000
         25%          16.282000
         50%          17.238000
         75%          18.212000
         max          19.995000
         Name: age, dtype: float64

```

```
In [53]: interests= teens.iloc[:,4:40]
         interests.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 36 columns):
basketball      30000 non-null int64
football        30000 non-null int64
soccer          30000 non-null int64
softball        30000 non-null int64
volleyball      30000 non-null int64
swimming        30000 non-null int64
cheerleading    30000 non-null int64
baseball        30000 non-null int64
tennis          30000 non-null int64
sports          30000 non-null int64
cute            30000 non-null int64
sex             30000 non-null int64
sexy            30000 non-null int64
hot             30000 non-null int64
kissed          30000 non-null int64
dance           30000 non-null int64
band            30000 non-null int64
marching        30000 non-null int64
music           30000 non-null int64
rock            30000 non-null int64

```



```

god          30000 non-null int64
church       30000 non-null int64
jesus        30000 non-null int64
bible        30000 non-null int64
hair         30000 non-null int64
dress        30000 non-null int64
blonde       30000 non-null int64
mall         30000 non-null int64
shopping     30000 non-null int64
clothes      30000 non-null int64
hollister    30000 non-null int64
abercrombie  30000 non-null int64
die          30000 non-null int64
death        30000 non-null int64
drunk        30000 non-null int64
drugs        30000 non-null int64
dtypes: int64(36)
memory usage: 8.2 MB

```

```

In [54]: from sklearn.preprocessing import StandardScaler
         scaler= StandardScaler()

```

```

In [55]: scaler.fit(interests)

```

```

Out[55]: StandardScaler(copy=True, with_mean=True, with_std=True)

```

```

In [56]: scaled_features= scaler.transform(interests)
         interests_z = pd.DataFrame(scaled_features, columns= interests.columns)
         interests_z.describe()

```

```

Out[56]:
          basketball      football      soccer      softball      volleyball \
count  3.000000e+04  3.000000e+04  3.000000e+04  3.000000e+04  3.000000e+04
mean   5.494864e-17 -4.547474e-17  1.515825e-17  3.031649e-17  3.789561e-18
std    1.000017e+00  1.000017e+00  1.000017e+00  1.000017e+00  1.000017e+00
min    -3.322173e-01 -3.576974e-01 -2.428741e-01 -2.179278e-01 -2.236696e-01
25%    -3.322173e-01 -3.576974e-01 -2.428741e-01 -2.179278e-01 -2.236696e-01
50%    -3.322173e-01 -3.576974e-01 -2.428741e-01 -2.179278e-01 -2.236696e-01
75%    -3.322173e-01 -3.576974e-01 -2.428741e-01 -2.179278e-01 -2.236696e-01
max     2.949277e+01  2.090850e+01  2.919421e+01  2.276453e+01  2.165366e+01

          swimming  cheerleading      baseball      tennis      sports \
count  3.000000e+04  3.000000e+04  3.000000e+04  3.000000e+04  3.000000e+04
mean   1.894781e-17 -3.789561e-17  3.789561e-18 -2.889540e-17 -2.842171e-17
std    1.000017e+00  1.000017e+00  1.000017e+00  1.000017e+00  1.000017e+00
min    -2.599706e-01 -2.073271e-01 -2.011306e-01 -1.689389e-01 -2.971234e-01
25%    -2.599706e-01 -2.073271e-01 -2.011306e-01 -1.689389e-01 -2.971234e-01
50%    -2.599706e-01 -2.073271e-01 -2.011306e-01 -1.689389e-01 -2.971234e-01
75%    -2.599706e-01 -2.073271e-01 -2.011306e-01 -1.689389e-01 -2.971234e-01

```

max	5.970348e+01	1.729137e+01	3.046682e+01	2.884728e+01	2.517666e+01
-----	--------------	--------------	--------------	--------------	--------------

	...	blonde	mall	shopping	clothes \
count	...	3.000000e+04	3.000000e+04	3.000000e+04	3.000000e+04
mean	...	-8.289665e-18	-1.136868e-17	5.873820e-17	1.515825e-17
std	...	1.000017e+00	1.000017e+00	1.000017e+00	1.000017e+00
min	...	-5.093652e-02	-3.699147e-01	-4.873142e-01	-3.141979e-01
25%	...	-5.093652e-02	-3.699147e-01	-4.873142e-01	-3.141979e-01
50%	...	-5.093652e-02	-3.699147e-01	-4.873142e-01	-3.141979e-01
75%	...	-5.093652e-02	-3.699147e-01	8.931794e-01	-3.141979e-01
max	...	1.683073e+02	1.687776e+01	1.469812e+01	1.661229e+01

		hollister	abercrombie	die	death	drunk \
count	3.000000e+04	3.000000e+04	3.000000e+04	3.000000e+04	3.000000e+04	
mean	5.494864e-17	1.136868e-17	-9.687066e-17	-1.610564e-17	-1.515825e-17	
std	1.000017e+00	1.000017e+00	1.000017e+00	1.000017e+00	1.000017e+00	
min	-2.014763e-01	-1.830317e-01	-2.947932e-01	-2.615302e-01	-2.204026e-01	
25%	-2.014763e-01	-1.830317e-01	-2.947932e-01	-2.615302e-01	-2.204026e-01	
50%	-2.014763e-01	-1.830317e-01	-2.947932e-01	-2.615302e-01	-2.204026e-01	
75%	-2.014763e-01	-1.830317e-01	-2.947932e-01	-2.615302e-01	-2.204026e-01	
max	2.575205e+01	2.843431e+01	3.493308e+01	3.179061e+01	1.982379e+01	

	drugs
count	3.000000e+04
mean	1.752672e-17
std	1.000017e+00
min	-1.749076e-01
25%	-1.749076e-01
50%	-1.749076e-01
75%	-1.749076e-01
max	4.613268e+01

[8 rows x 36 columns]

```
In [57]: from sklearn.cluster import KMeans
```

```
In [58]: kmeans = KMeans(n_clusters=5)
```

```
In [59]: kmeans.fit(interests_z)
```

```
Out[59]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
                n_clusters=5, n_init=10, n_jobs=1, precompute_distances='auto',
                random_state=None, tol=0.0001, verbose=0)
```

```
In [60]: kmeans.labels_
```

```
Out[60]: array([0, 2, 0, ..., 0, 0, 0], dtype=int32)
```

```
In [61]: labels = pd.DataFrame(kmeans.labels_)
```

```
In [62]: teens_labels = pd.concat([teens,labels], axis=1)
```

```
In [63]: teens_labels.rename(columns={0: 'labels'}, inplace=True)
```

```
In [64]: teens_labels.head()
```

```
Out[64]:
```

	gradyear	gender	age	friends	basketball	football	soccer	softball	\
0	2006	M	18.982	7	0	0	0	0	
1	2006	F	18.801	0	0	1	0	0	
2	2006	M	18.335	69	0	1	0	0	
3	2006	F	18.875	0	0	0	0	0	
4	2006	NaN	18.995	10	0	0	0	0	

	volleyball	swimming	...	clothes	hollister	abercrombie	die	death	\
0	0	0	...	0	0	0	0	0	
1	0	0	...	0	0	0	0	0	
2	0	0	...	0	0	0	0	1	
3	0	0	...	0	0	0	0	0	
4	0	0	...	0	0	0	0	0	

	drunk	drugs	F	M	labels
0	0	0	0	1	0
1	0	0	1	0	2
2	0	0	0	1	0
3	0	0	1	0	0
4	1	1	0	0	1

[5 rows x 43 columns]

```
In [65]: teens_labels['labels'].value_counts()
```

```
Out[65]: 0    21530
         2     4201
         3     2635
         1     1035
         4       599
         Name: labels, dtype: int64
```

```
In [66]: AgeMeansByLabels = teens_labels['age'].groupby(teens_labels['labels']).mean()
         print(AgeMeansByLabels)
```

```
labels
0    17.300643
1    17.098047
2    17.050091
3    17.040383
4    17.381677
         Name: age, dtype: float64
```

```
In [67]: FemaleMeansByLabels = teens_labels['F'].groupby(teens_labels['labels']).mean()
        print(FemaleMeansByLabels)
```

```
labels
0    0.706781
1    0.801932
2    0.887170
3    0.700949
4    0.722871
Name: F, dtype: float64
```

```
In [68]: FriendsMeanByLabels = teens_labels['friends'].groupby(teens_labels['labels']).mean()
        print(FriendsMeanByLabels)
```

```
labels
0    27.772782
1    30.727536
2    38.387527
3    35.909298
4    32.964942
Name: friends, dtype: float64
```

```
In [69]: teens.columns
```

```
Out[69]: Index([u'gradyear', u'gender', u'age', u'friends', u'basketball', u'football',
               u'soccer', u'softball', u'volleyball', u'swimming', u'cheerleading',
               u'baseball', u'tennis', u'sports', u'cute', u'sex', u'sexy', u'hot',
               u'kissed', u'dance', u'band', u'marching', u'music', u'rock', u'god',
               u'church', u'jesus', u'bible', u'hair', u'dress', u'blonde', u'mall',
               u'shopping', u'clothes', u'hollister', u'abercrombie', u'die', u'death',
               u'drunk', u'drugs', u'F', u'M'],
              dtype='object')
```

```
In [77]: from scipy.spatial.distance import cdist
        distortions = []
        for i in range(1,40):
            kmeans = KMeans(n_clusters=i)
            kmeans.fit(interests_z)
            distortions.append(sum(np.min(cdist(interests_z, kmeans.cluster_centers_, 'euclidean')
```

```
In [80]: plt.plot(range(1,40), distortions, 'bx-')
        plt.xlabel('k')
        plt.ylabel('Distortion')
        plt.title('The Elbow Method showing the optimal k')
        plt.show()
```

